61A Exam Data Analysis: Measure of Academic Dishonesty

Kavi Gupta

January 9, 2017

1   General Strategy

The general strategy is to look at similarities between scores of students sitting next to each other against the similarities between scores of students sitting not next to each other.

The result can be then used to estimate the number of students committing academic dishonesty using a model.

2   Confounding Factors

There are many potential confounding factors, which we control for in different ways.

2.1   Grader Differences

Some graders grade differently from each other. Taking Midterm 1 problem 1.3, we can see that graders 5 and 8 had fairly typical behavior in terms of average scores given. However, we can also see that they have a very different profile of rubric items. In this case, what happened was that these two graders were giving rubric items 2, 3, and 4, which corresponded to each of the three parts of the problem being correct, rather than giving rubric item 1, which corresponded to the entire problem being correct. While this is irrelevant to the student, it does confound analyses.

Since grading ranges tend to be at least somewhat associated with location, this might lead to students near each other having more similar profiles, artificially. We control for graders by subtracting out from each student's individual problem scores and rubric items the mean given by a grader.

Any student who had at least one problem that was either graded by a grader who had graded fewer than ten problems (not enough of a track record) or who had an abnormal pattern of rubric item grades[1].

---

[1]Unusualness of a grader is defined as

$$u = \sum_k \left| \frac{x_k - \mu_k}{\sigma_k} \right|$$

where x_k is the mean for the given grader of the kth rubric item, and μ_k, σ_k are the mean and standard deviation for that rubric item in general

2.2 Sequencing

The Gambler's Fallacy could possibly lead to a negative correlation between consecutively graded exams. To control for this, we simply ignore pairs of exams that are graded near each other in time.

2.3 Room

Individual rooms might have different rubric profiles, due to, for example, different TAs in different rooms. It looks like no particular room has that much of a problem, so for now rooms are not controlled for except to not take pairs from different rooms (there is a small amount of variation).

2.4 Front of Room / Back of Room

Sections within rooms appear to have little effect on exam profiles (at least within the margin of error), and definitely nothing systematic.

2.5 Aisle / Middle

No analysis on this factor as of yet. Trying to find aisle locations is impossible from seating chart data; would require some data gathering.

2.6 Seating Chart Inaccuracies

Not yet fixed. Have the data, but would require a lot of data input because OCR is not particularly effective at recognizing handwriting.

3 Measuring Similarity

3.1 Correlation

Using correlation between normalized scores as a measure of similarity, we can prove that there is a difference between the similarities in the set of adjacent pairs of students and the set of non-adjacent pairs of students (see controls above for more details).

Overall, a permutation test shows that the difference is in fact significant:

3.2 Absolute Difference in Overall Exam Score

An alternate means of similarity is far simpler. We look at the average difference between the scores in the two exams. It turns out that despite its unsophistication, this test is also able to see past the noise in the data.


In this case, the difference is the opposite as above, because a higher correlation corresponds to a smaller absolute difference.

4 Model For Academic Dishonesty

4.1 Score-Independent, Binary Cheater, Random Cheat-Selection, Model

Assume that some fixed fraction c of the students cheat by copying exactly k randomly selected points worth of material from a person sitting next to them (also selected randomly). Assume that non-cheating students get each point independently with probability p. Further assume that cheating students are located randomly.

Using known exam data, we can approximate p simply by taking the average score divided by the total number of points.

Score Independent Model Scores vs. Real Scores
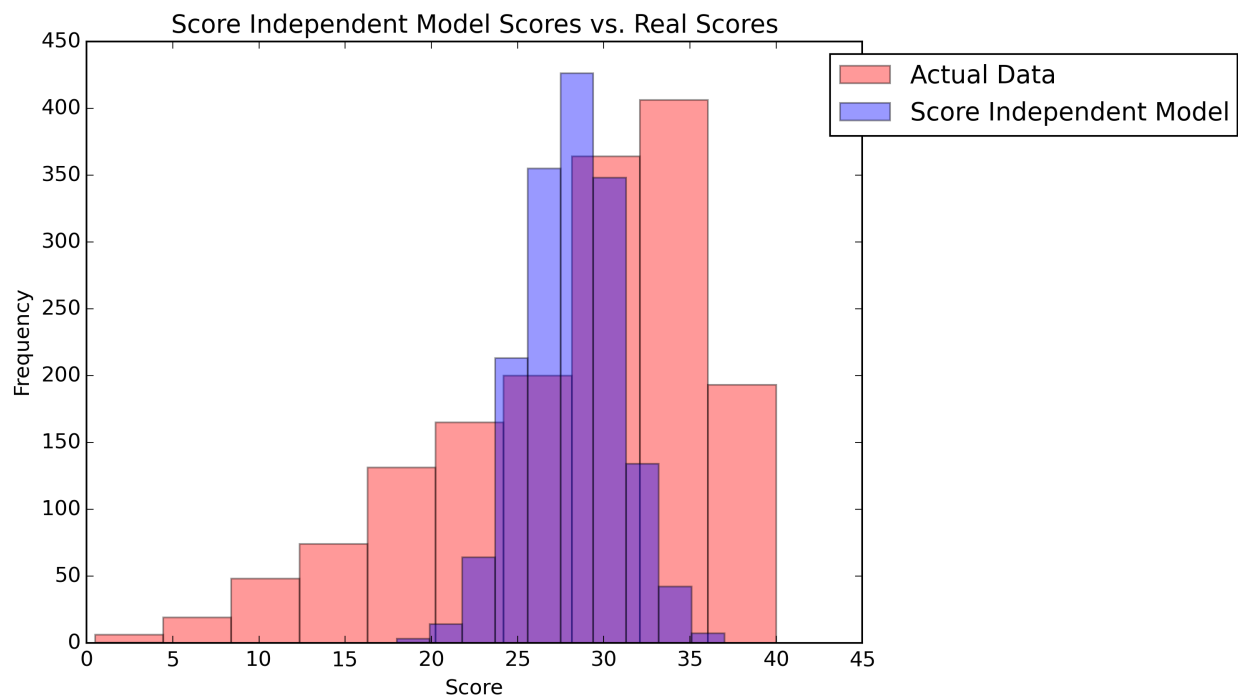
Unfortunately, any model generated using these does not fit the data. This is because the variance in this model is far smaller than the variance in the actual data.

4.2   Chunked Score-Independent, Binary Cheater, Random Cheat-Selection Model

This model is equivalent to the one above, but uses point chunks representing each question, which are assumed to be independent and take the sizes of each of the questions. A certain number of questions rather than points are copied. Each question is assigned a score based on a normal distribution from the